



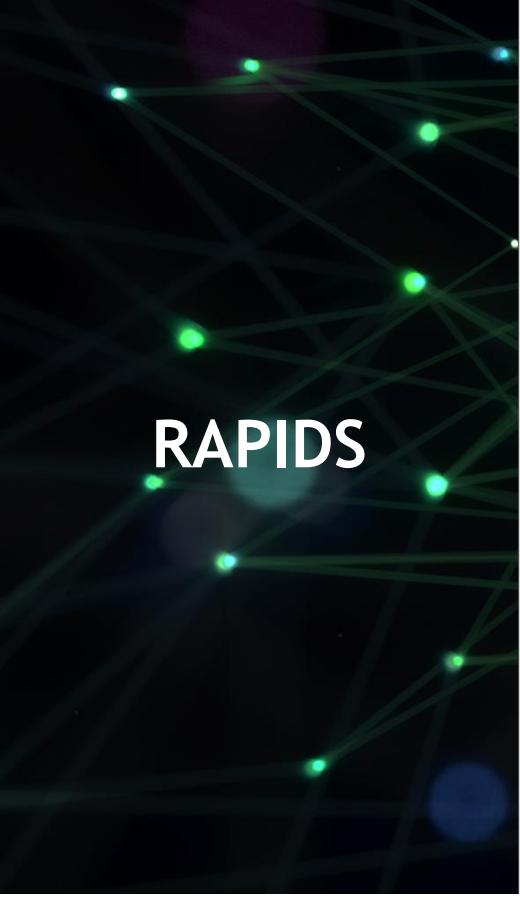
RAPIDS

GPU POWERED MACHINE LEARNING

Miguel Martínez

WHAT IS RAPIDS





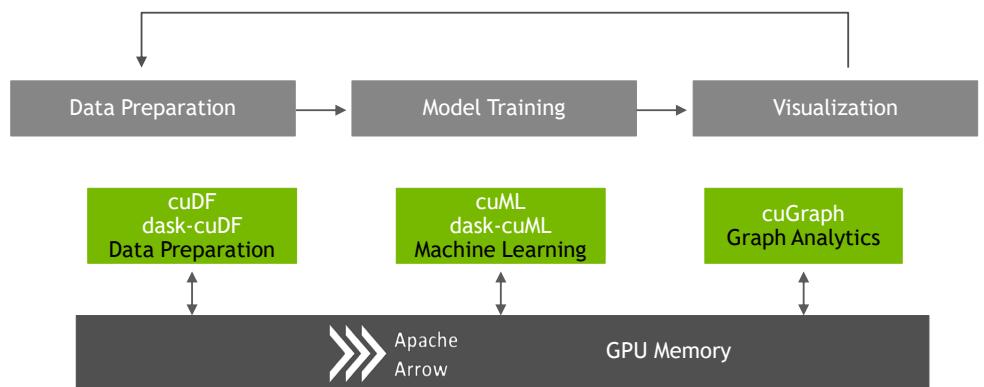
RAPIDS

RAPIDS

GPU Accelerated End-to-End Data Science

RAPIDS is a set of open source libraries for GPU accelerating **data preparation** and **machine learning**.

OSS website: rapids.ai





RAPIDS

RAPIDS LIBRARIES

cuDF

- GPU-accelerated lightweight in-GPU memory database used for data preparation
- Accelerates loading, filtering, and manipulation of data for model training data preparation
- Python drop-in Pandas replacement built on CUDA C++

cuML

- GPU accelerated traditional machine learning libraries
- XGBoost, PCA, Kalman, K-means, k-NN, DBScan, tSVD ...

cuGRAPH

- Collection of graph analytics libraries.

HOW TO SETUP AND START USING RAPIDS



HOW? DOWNLOAD AND DEPLOY

Source available on GitHub | Container available on NGC and Docker Hub | Conda and PIP

<https://github.com/rapidsai> <https://ngc.nvidia.com> <https://pypi.org/project/rapidsai>

<https://hub.docker.com/u/rapidsai> <https://anaconda.org/rapidsai>

<https://pypi.org/project/cudf/> <https://pypi.org/project/cuml/>

GitHub



NGC



docker

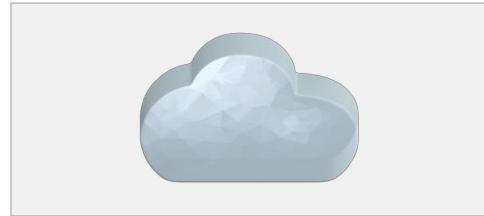
CONDA



CONDA



On-premises



Cloud

*Pascal GPU architecture or better
CUDA 9.2 or 10.0
Ubuntu 16.04 or 18.04*

RUNNING RAPIDS CONTAINER IN THE CLOUD

A step-by-step installation guide (MS Azure)

1. Create a *NC6s_v2* virtual machine instance on *Microsoft Azure Portal* using *NVIDIA GPU Cloud Image for Deep Learning and HPC* as image.

2. Start the virtual machine.

3. Connect to the virtual machine using the following command:

```
$ ssh -L 8080:localhost:8888 \
-L 8787:localhost:8787 \
username@public_ip_address
```

4. Pull the *RAPIDS container* from *NGC*. Run it.

```
$ docker pull nvcr.io/nvidia/rapidsai/rapidsai:cuda10.0-runtime-ubuntu18.04
$ docker run --runtime=nvidia \
--rm -it \
-p 8888:8888 \
-p 8787:8787 \
-p 8786:8786 \
nvcr.io/nvidia/rapidsai/rapidsai:cuda10.0-runtime-ubuntu18.04
```

5. Run JupyterLab:

```
(rapids)$ bash /rapids/notebooks/utils/start-jupyter.sh
```

6. Open your browser, and navigate to <http://localhost:8080>.

7. Navigate to:

- *cuml* folder for cuML IPython examples.
- *mortgage* folder for XGBoost IPython examples.

8. Enjoy!

RUNNING RAPIDS CONTAINER IN THE CLOUD

A step-by-step installation guide (AWS)

1. Create a *p3.8xlarge* machine instance on Amazon Web Services using *NVIDIA Volta Deep Learning AMI* as image.
2. Start the virtual machine.
3. Connect to the virtual machine using the following command:

```
$ ssh -L 8080:localhost:8888 \
-L 8787:localhost:8787 \
ubuntu@public_ip_address
```
4. Pull the *RAPIDS container* from NGC. Run it.

```
$ docker pull nvcr.io/nvidia/rapidsai/rapidsai:cuda10.0-runtime-ubuntu18.04
$ docker run --runtime=nvidia \
--rm -it \
-p 8888:8888 \
-p 8787:8787 \
-p 8786:8786 \
nvcr.io/nvidia/rapidsai/rapidsai:cuda10.0-runtime-ubuntu18.04
```

5. Run JupyterLab:
(rapids)\$ bash /rapids/notebooks/utils/start-jupyter.sh
6. Open your browser, and navigate to <http://localhost:8080>.
7. Navigate to:
 - *cuml* folder for cuML IPython examples.
 - *mortgage* folder for XGBoost IPython examples.
8. Enjoy!

HOW TO PORT EXISTING CODE

Principal Component Analysis (PCA)

Before... ...Now!

CPU vs GPU

PORTING EXISTING CODE

PCA

Training results:

- CPU: 57.1 seconds
- GPU: 4.28 seconds

System: AWS p3.8xlarge
CPUs: Intel(R) Xeon(R) CPU E5-2686 v4 @ 2.30GHz, 32 vCPU cores, 244 GB RAM
GPU: Tesla V100 SXM2 16GB
Dataset: <https://github.com/rapidsai/cuml/tree/master/python/notebooks/data>

Specific: Import CPU algorithm

```
[1]: from sklearn.decomposition import PCA
```

Common: Helper functions

```
[2]: # Timer, Load_data...
from helper import *
```

Common: Data loading and algo params

```
[3]: # Data Loading
nrows = 2**22
ncols = 400

X = load_data(nrows, ncols)
print("data", X.shape)

# Algorithm parameters
n_components = 8
whiten = False
random_state = 42
svd_solver = "full"

use mortgage data
data (4194304, 400)
```

Specific: Import GPU algorithm

```
[1]: from cuml import PCA
```

Common: Helper functions

```
[2]: # Timer, Load_data...
from helper import *
```

Common: Data loading and algo params

```
[3]: # Data Loading
nrows = 2**22
ncols = 400

X = load_data(nrows, ncols)
print("data", X.shape)

# Algorithm parameters
n_components = 10
whiten = False
random_state = 42
svd_solver = "full"

use mortgage data
data (4194304, 400)
```

Specific: DataFrame from Pandas to cuDF

```
[4]: %%time
import cudf
X = cudf.DataFrame.from_pandas(X)

CPU times: user 4.46 s, sys: 4.68 s, total: 9.14 s
Wall time: 9.36 s
```

Common: Training

```
[4]: %%time
pca = PCA(n_components=n_components, svd_solver=svd_solver,
           whiten=whiten, random_state=random_state)
_ = pca.fit_transform(X)
```

CPU times: user 9min 19s, sys: 2min 12s, total: 11min 32s

Wall time: 57.1 s

```
[5]: %%time
pca = PCA(n_components=n_components, svd_solver=svd_solver,
           whiten=whiten, random_state=random_state)
_ = pca.fit_transform(X)
```

CPU times: user 1.94 s, sys: 512 ms, total: 2.45 s

Wall time: 4.28 s

k-Nearest Neighbors (KNN)

Before...

...Now!

Specific: Import CPU algorithm

```
[1]: from sklearn.neighbors import KDTree as KNN
```

Common: Helper functions

```
[2]: # Timer, Load_data...
from helper import *
```

Common: Data loading and algo params

```
[3]: # Data Loading
nrows = 2**17
ncols = 40

X = load_data(nrows, ncols)
print('data', X.shape)

# Algorithm parameters
n_neighbors = 10

use mortgage data
data (131072, 40)
```

Specific: Import GPU algorithm

```
[1]: from cuml import KNN
```

Common: Helper functions

```
[2]: # Timer, Load_data...
from helper import *
```

Common: Data loading and algo params

```
[3]: # Data Loading
nrows = 2**17
ncols = 40

X = load_data(nrows, ncols)
print('data', X.shape)

# Algorithm parameters
n_neighbors = 10

use mortgage data
data (131072, 40)
```

Specific: DataFrame from Pandas to cuDF

```
[4]: %time
import cudf
X = cudf.DataFrame.from_pandas(X)

CPU times: user 3 s, sys: 552 ms, total: 3.56 s
Wall time: 839 ms
```

Specific: Training

```
[4]: %time
knn = KNN(X)
_ = knn.fit(X, n_neighbors)

CPU times: user 9min 2s, sys: 272 ms, total: 9min 2s
Wall time: 8min 59s
```

Specific: Training

```
[5]: %time
knn = KNN(n_gpus=1)
knn.fit(X)
_ = knn.query(X, n_neighbors)

CPU times: user 692 ms, sys: 428 ms, total: 1.12 s
Wall time: 1.12 s
```

CPU vs GPU PORTING EXISTING CODE KNN

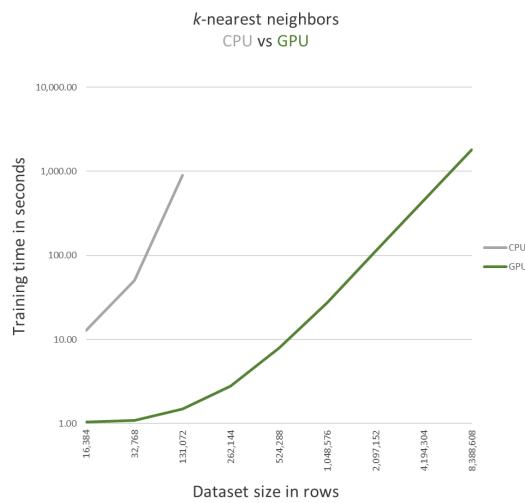
Training results:

- CPU: ~9 minutes
- GPU: 1.12 seconds

System: AWS p3.8xlarge
CPUs: Intel(R) Xeon(R) CPU E5-2686 v4 @ 2.30GHz, 32 vCPU cores, 244 GB RAM
GPU: Tesla V100 SXM2 16GB
Dataset: <https://github.com/rapidsai/cuml/tree/master/python/notebooks/data>

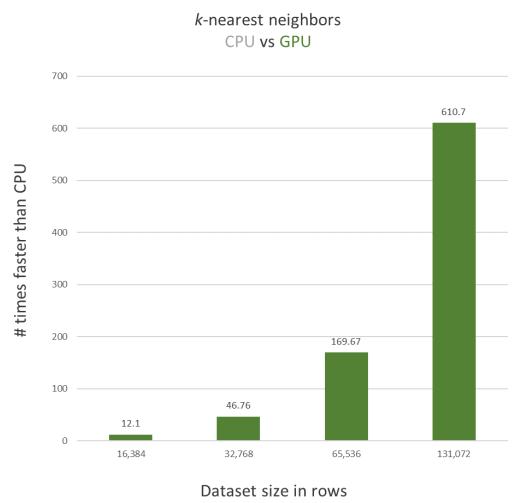
TRAINING TIME COMPARISON

CPU vs GPU



Dataset size trained in 15 minutes.
CPU: ~130.000 rows.
GPU: ~5.900.000 rows.

Specs	NC6s_vs
Cores (Broadwell 2.6Ghz)	6
GPU	1 x P100
Memory	112 GB
Local Disk	-700 GB SSD
Network	Azure Network



The bigger the dataset is, the higher the training performance difference is between CPU and GPU.

WHAT IS XGBOOST

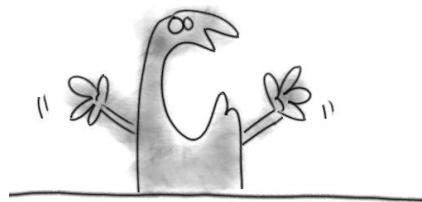
XGBOOST

Definition

“

XGBoost is an implementation of gradient boosted decision trees designed for speed and performance.

What?!!



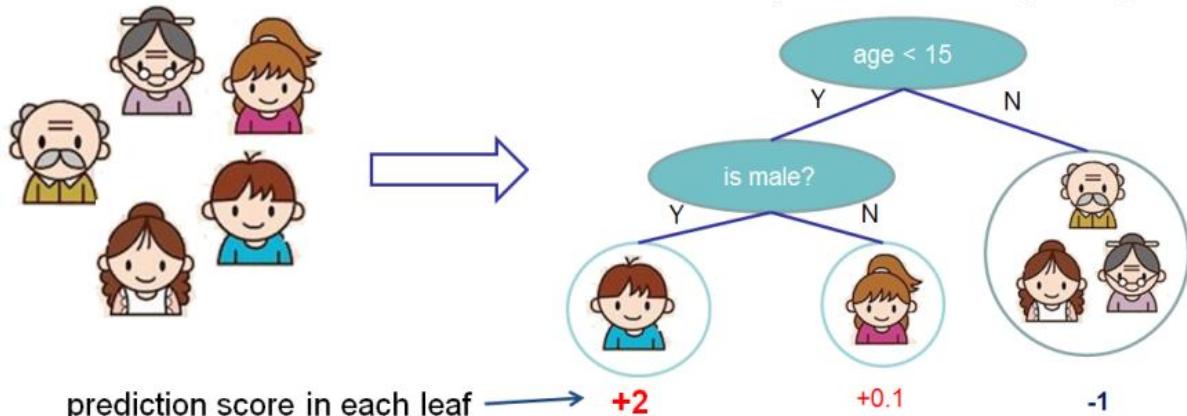
It is a powerful tool for solving classification and regression problems in a supervised learning setting.

PREDICT: WHO ENJOYS COMPUTER GAMES

Example of Decision Tree

Input: age, gender, occupation, ...

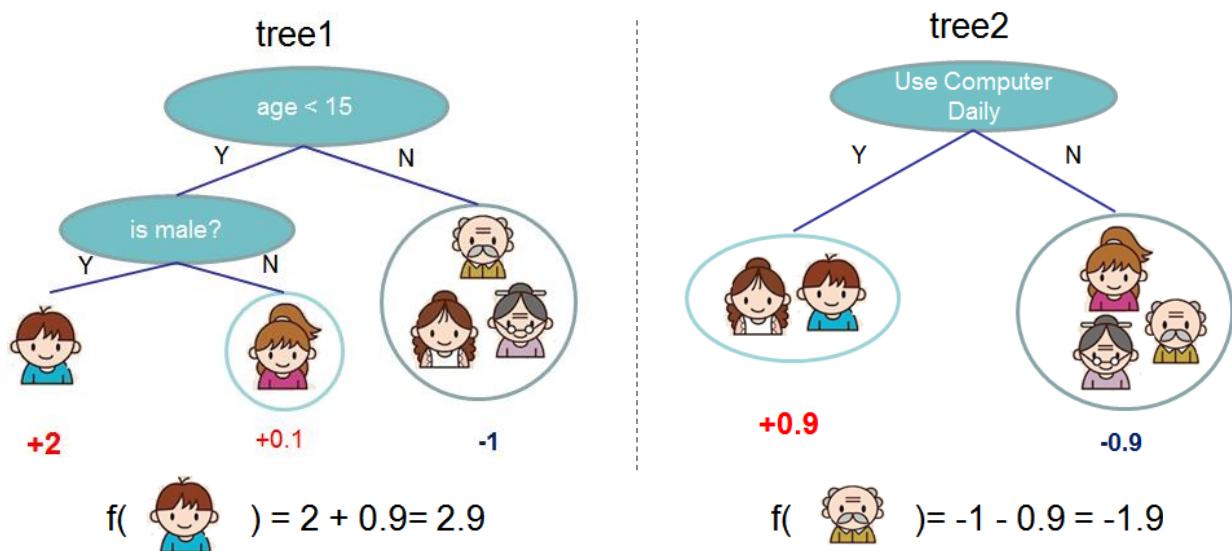
Does the person like computer games



Source: <https://goo.gl/C6WKif>

COMBINE TREES FOR STRONGER PREDICTIONS

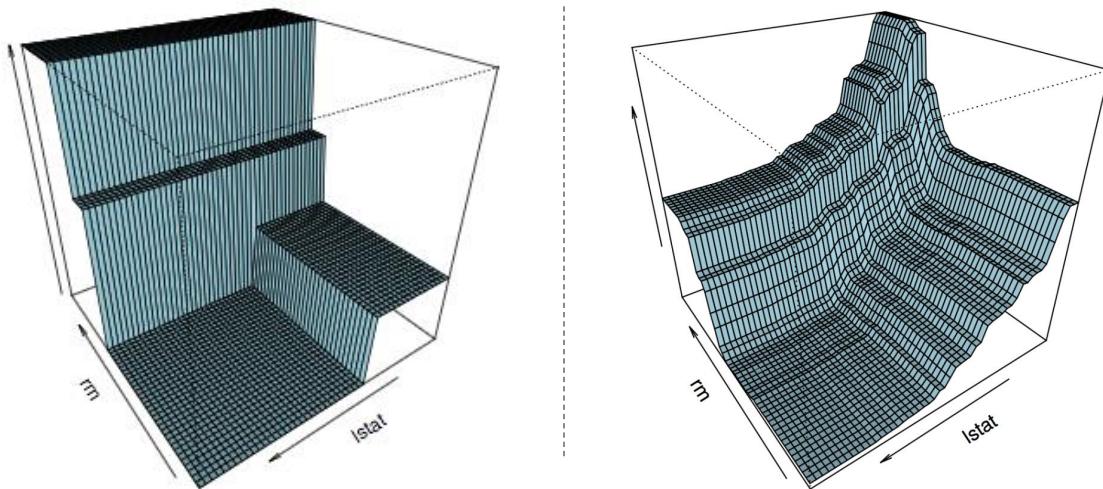
Example of Using Ensembled Decision Trees



Source: <https://goo.gl/C6WKif>

TRAINED MODELS VISUALIZATION

Single Decision Tree vs Ensembled Decision Trees



Models fit to the *Boston Housing* Dataset.

Source: <https://goo.gl/GWNdEm>

WHY XGBoost



A STRONG HISTORY OF SUCCESS

On a Wide Range of Problems

Winner of Caterpillar Kaggle Contest 2015 

- Machinery component pricing

Winner of CERN Large Hadron Collider Kaggle Contest 2015 

- Classification of rare particle decay phenomena

Winner of KDD Cup 2016 

- Research institutions' impact on the acceptance of submitted academic papers

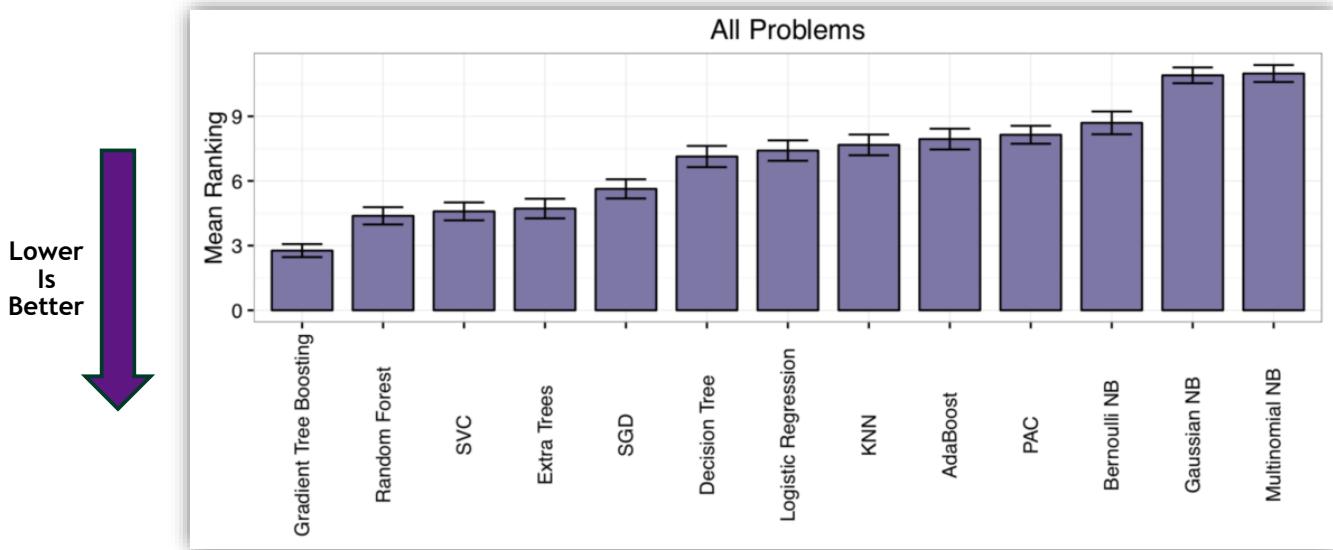
Winner of ACM RecSys Challenge 2017 

- Job posting recommendation



WHICH ML ALGORITHM PERFORMS BEST

Average rank across 165 ML datasets



WHY XGBOOST + RAPIDS



WHY RAPIDS WITH XGBOOST

Multi-GPU, Multi-Node, Scalability

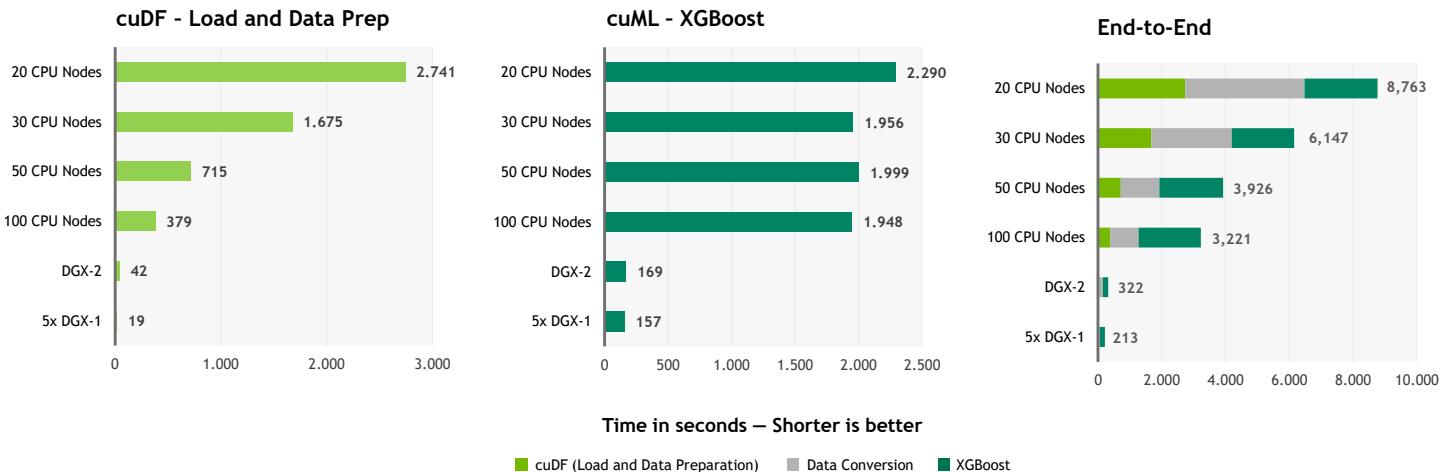
► XGBoost:

- ▶ Algorithm tuned for eXtreme performance and high efficiency
- ▶ Multi-GPU and Multi-Node Support

► RAPIDS:

- ▶ End-to-end data science & analytics pipeline entirely on GPU
- ▶ User-friendly Python interfaces
- ▶ Faster results helps hyperparameter tuning
- ▶ Relies on CUDA primitives, exposes parallelism and high-memory bandwidth

BENCHMARKS



Benchmark

200GB CSV dataset; Data preparation includes joins, variable transformations.

CPU Cluster Configuration

CPU nodes (61 GiB of memory, 8 vCPUs, 64-bit platform), Apache Spark

DGX Cluster Configuration

5x DGX-1 on InfiniBand network

cuML ROADMAP



cuML Algorithms	Available Now	Q2-2019
XGBoost GBDT	MGMN	
XGBoost Random Forest		MGMN
K-Means Clustering	MG	
K-Nearest Neighbors (KNN)	MG	
Principal Component Analysis (PCA)	SG	
Density-based Spatial Clustering of Applications with Noise (DBSCAN)	SG	
Truncated Singular Value Decomposition (tSVD)	SG	
Uniform Manifold Aproximation and Projection (UMAP)	SG	MG
Kalman Filters (KF)	SG	
Ordinary Least Squares Linear Regression (OLS)	SG	
Stochastic Gradient Descent (SGD)	SG	
Generalized Linear Model, including Logistic (GLM)		MG
Time Series (Holts-Winters)		SG
Autoregressive Integrated Moving Average (ARIMA)		SG

SG
Single GPU

MG
Multi-GPU

MGMN
Multi-GPU Multi-Node

Last updated 29.03.19



LEARN MORE ABOUT RAPIDS

HOME

GETTING STARTED

COMMUNITY

GITHUB

BLOG

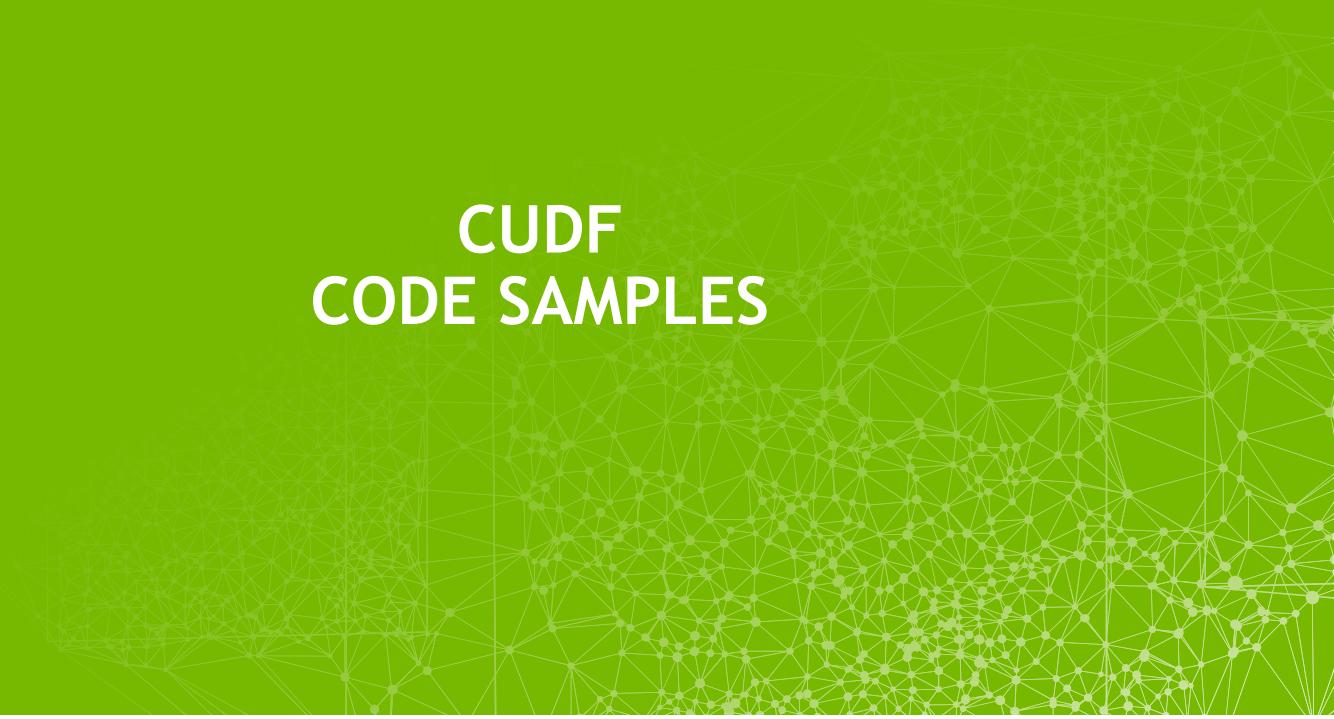
RAPIDS

Open GPU Data Science

GET STARTED

<https://rapids.ai>

CUDF CODE SAMPLES



LOADING DATA INTO A GPU DATAFRAME

Create an empty DataFrame, and add a column.

```
[1]: import cudf  
  
gdf = cudf.DataFrame()  
gdf['my_column'] = [6, 7, 8]  
print(gdf)
```

my_column
0 6
1 7
2 8

Create a DataFrame with two columns.

```
[2]: import cudf  
  
gdf = cudf.DataFrame({'a': [3, 4, 5], 'b': [6, 7, 9]})  
  
print(gdf)
```

b	a
0 6	3
1 7	4
2 9	5

Load a CSV file into a GPU DataFrame.

```
[3]: import cudf  
  
path = './apartments.csv'  
  
names = ['city', 'zipcode', 'price_per_m2', 'year_built',  
         'population', 'median_income', 'date']  
  
# Note: dtype detection is not yet supported.  
dtypes = ['category', 'int64', 'float64', 'float64',  
          'int64', 'int64', 'date']  
  
gdf = cudf.read_csv(path, names=names, dtype=dtypes, delimiter=';',  
                    skiprows=1, skipfooter=1)
```

Use Pandas to load a CSV file, and copy its content into a GPU DataFrame.

```
[4]: import pandas as pd  
import cudf  
  
# Load a CSV file using pandas.  
pdf = pd.read_csv(path, delimiter=',')  
  
# Convert data types to ones supported by cudf.  
pdf['city'] = pdf['city'].astype('category')  
pdf['date'] = pdf['date'].astype("datetime64[ms]")  
  
# Create a cudf dataframe from a pandas dataframe.  
gdf = cudf.DataFrame.from_pandas(pdf)
```

WORKING WITH GPU DATAFRAMES

Return the first three rows as a new DataFrame.

```
[5]: print(gdf.head(3))
```

	city	zipcode	price_per_m2	year_built	population	median_income	date
0	espoo	2100	5444.022222	1985	4332	26167	2018-09-06T00:00:00.000
1	espoo	2130	3768.0	1972	5983	29579	2018-08-20T00:00:00.000
2	espoo	2140	2770.0	1977	3689	29447	2018-12-19T00:00:00.000

Row slicing with column selection.

```
[6]: print(gdf.loc[2:5, ['zipcode', 'year_built']])
```

	zipcode	year_built
2	2140	1977
3	2160	1990
4	2170	1972
5	2180	1986

Find the mean and standard deviation of a column.

```
[7]: print(gdf['population'].mean())
print(gdf['population'].std())
```

8014.397849462365
4373.122998945762

Count number of occurrences per value, and number of unique values.

```
[8]: print(gdf['city'].value_counts())
print(gdf['city']).unique_count() # nunique() in pandas.
```

city	count
helsinki	65
espoo	28
2	2

Change the data type of a column.

```
[9]: import numpy as np

print('Median income dtype used to be:', gdf['median_income'].dtype)
gdf['median_income'] = gdf['median_income'].astype(np.float64)
print('Median income dtype is now:', gdf['median_income'].dtype)
```

Median income dtype used to be: int64
Median income dtype is now: float64

Transform column values with a custom function.

```
[10]: def double_income(median_income):
        return 2*median_income

gdf['median_income'] = gdf['median_income'].applymap(double_income)

print(gdf.head(2))
```

	city	zipcode	price_per_m2	year_built	population	median_income	date
0	espoo	2100	5444.022222	1985	4332	52334.0	2018-09-06T00:00:00.000
1	espoo	2130	3768.0	1972	5983	59158.0	2018-08-20T00:00:00.000

QUERY, SORT, GROUP, JOIN, MERGE, ONE-HOT ENCODING

Query the columns of a DataFrame with a boolean expression.

```
[11]: query = gdf.query("year_built < 1930")
print(query.head(3))
```

```
city zipcode      price_per_m2 year_built population median_income          date
30 helsinki     130        7916.0    1911       1536      56226.0 2019-02-17T00:00:00.000
31 helsinki     140  7416.905659999999  1925       7817      55194.0 2018-10-09T00:00:00.000
32 helsinki     150  7727.571429000005  1907       9299      49734.0 2018-09-29T00:00:00.000
```

Return the first 'n' rows ordered by 'columns' in ascending order.

```
[13]: three_smallest = gdf.nsmallest(n=3, columns=['population'])
print(three_smallest)
```

```
city zipcode price_per_m2 year_built population median_income          date
41 helsinki     310        3971.0    1972       896      46688.0 2018-10-05T00:00:00.000
30 helsinki     130        7916.0    1911       1536      56220.0 2019-02-17T00:00:00.000
44 helsinki     340  4497.333333  1973       1654      64768.0 2018-11-20T00:00:00.000
```

Join columns with other DataFrame on index.

```
[15]: left = grouped
right = cudf.DataFrame({'zipcode': [28, 65], 'feature1': [1,2]})

# join() uses the index.
join_left = left.set_index('count_zipcode')
join_right = right.set_index('zipcode')

# Different join styles are supported.
joined = join_left.join(join_right, how='right')
```

Sort a column by its values.

```
[12]: gdf = gdf.sort_values(by='population', ascending=False)
print(gdf.head(3))
```

```
city zipcode      price_per_m2 year_built population median_income          date
89 helsinki     940        1982.028571  1967       25817      38172.0 2019-02-18T00:00:00.000
  8 espoo        2230        4035.075   1992       20397      46148.0 2018-12-09T00:00:00.000
58 helsinki     530        5090.853659  1944       18663      42582.0 2018-10-03T00:00:00.000
```

Group by column with aggregate function.

```
[14]: # Differences to pandas:
# - aggregated column names are prefixed with the
#   aggregated function name.
# - 'city' becomes index in pandas but not in cudf.
grouped = gdf.groupby(['city']).agg({'zipcode': 'count'})
```

Merge two DataFrames.

```
[16]: # Only inner join is supported currently.
merged = left.merge(right, on=['zipcode'])
```

One-hot encoding.

```
[17]: gdf['city_codes'] = gdf.city.cat.codes
codes = gdf.city_codes.unique()

# get_dummies() in pandas.
encoded = gdf.one_hot_encoding(column='city_codes', cats=codes,
                                prefix='city_codes_dummy', dtype='int8')
```

SUMMARY



RAPIDS

GPU Accelerated Data Science

RAPIDS is a set of open source libraries for GPU accelerating **data preparation** and **machine learning**.

Visit www.rapids.ai



ONE MORE THING

FIND A NEW ARGUMENT

THE #1 DATA SCIENTIST EXCUSE
FOR LEGITIMATELY SLACKING OFF:

"MY MODEL'S TRAINING . "



MESSAGE TO
DATA SCIENTISTS

